

EDA_ENZYMES_group9

December 6, 2020

1 Data Mining Programming Assignment

2 Task1 - Exploratory Data Analysis

2.0.1 Group 9 - Nikola Vinko / Peter Hunyadi

We assume that you have the generated Gram matrices and sparse vectors in the directories

kernels/node_labels

kernels/without_labels

And have the HiSC repository in: `> ./HiSC`

If not you can download them [here](https://github.com/maxfat/data_mining_group9.git). Repository: https://github.com/maxfat/data_mining_group9.git

```
[2]: # Load packages
import os
import numpy as np
from auxiliarymethods import datasets as dp
from auxiliarymethods.reader import tud_to_networkx
import pandas as pd
from sklearn.decomposition import KernelPCA, TruncatedSVD, PCA
from matplotlib import pyplot as plt
from nrkmeans import NrKmeans
from sklearn.cluster import SpectralClustering, KMeans, AgglomerativeClustering
from sklearn.metrics import normalized_mutual_info_score, pairwise_distances
import seaborn as sns
from scipy.sparse import load_npz
import auxiliarymethods.auxiliary_methods as aux
import networkx as nx
from sklearn.cluster import KMeans
from copy import deepcopy
from sklearn.manifold import TSNE
import umap
from sklearn.preprocessing import StandardScaler
from sklearn.manifold import TSNE
from sklearn.model_selection import train_test_split
from sklearn import preprocessing, svm, metrics
```

```

from scipy.sparse import csc_matrix
from sklearn.cluster import OPTICS, cluster_optics_dbscan
import matplotlib.gridspec as gridspec
import sys
sys.path.append(r'./HiSC/data_mining_group9/intermediate submission/') #may
from HiSC import process_csv, HiSC, reachability_plot, get_clusters

```

```

[4]: # functions
def load_csv(path):
    return np.loadtxt(path, delimiter=";")

def load_sparse(path):
    return load_npz(path)

def select_from_list(l, indices):
    return [l[i] for i in indices]

def visualize(G, color=None, figsize=(5,5)):
    plt.figure(figsize=figsize)
    plt.xticks([])
    plt.yticks([])
    nx.draw_networkx(G,
                     pos=nx.spring_layout(G, seed=42),
                     with_labels=True,
                     node_color=color,
                     cmap="Set2")
    plt.show();

def plot_dr(reduced_dim, title, labels):
    targets = [1, 2, 3, 4, 5, 6]
    colors = ['r', 'g', 'b', 'y', 'orange', "purple"]

    fig, ax = plt.subplots(figsize=(10,10))
    for color, target in zip(colors, targets):
        ind_to_keep = []
        for i in range(0, len(reduced_dim)):
            if labels[i] == target:
                ind_to_keep.append(i)

        x = reduced_dim[ind_to_keep, 0]
        y = reduced_dim[ind_to_keep, 1]
        ax.scatter(x, y, c=color, label=target, alpha = 0.8)
    ax.legend()
    plt.title(title)
    plt.show();

```

```
[5]: #main named variables
base_path = os.path.join("kernels", "node_labels")
dataset = "ENZYMES"
classes = dp.get_dataset(dataset)
G = tud_to_networkx(dataset)
print(f"Number of graphs in data set is {len(G)}")
print(f"Number of classes {len(set(classes.tolist()))}")
```

Number of graphs in data set is 600
Number of classes 6

3 First look of the data: Dimensionality Reduction on raw data - WL kernel

Let's see how our data looks like when unprocessed

```
[6]: # trying different iterations
dataset = "ENZYMES"
iterations = 5
base_path = os.path.join("kernels", "node_labels")
#base_path = os.path.join("kernels", "without_labels")

#-----
# Load graph representation from the kernel
#-----
#Gram Matrix for the Weisfeiler-Lehman subtree kernel
gram = load_csv(os.path.join(base_path, f"{dataset}_gram_matrix_wl{iterations}."
    ↪ "csv"))
gram = aux.normalize_gram_matrix(gram)

#Sparse Vectors for the Weisfeiler-Lehmann subtree kernel
vec = load_sparse(os.path.join(base_path, f"{dataset}_vectors_wl{iterations}."
    ↪ "npz"))
print(gram.shape, vec.shape)

# SVD
tsvd = TruncatedSVD(n_components=50)
reduced_tsvd = tsvd.fit_transform(vec)
plot_dr(reduced_tsvd, "TSVD", classes)

# Kernel PCA
kpca = KernelPCA(n_components=100, kernel="precomputed")
reduced_kpca = kpca.fit_transform(gram)
plot_dr(reduced_kpca, "Kernel PCA", classes)

# T-SNE
```


```


pca = TSNE(n_components=3)
reduced_tsne = pca.fit_transform(gram)
plot_dr(reduced_tsne, "T-SNE", classes)


# UMAP
reducer = umap.UMAP()
scaled_gram = StandardScaler().fit_transform(gram)
embedding = reducer.fit_transform(scaled_gram)
plot_dr(embedding, "UMAP", classes)


```

(600, 600) (600, 76881)









We can notice some outliers that occur in several vector representations, especially in kPCA. Let's visualize them to see if we can notice some patterns


```
[7]: # lets investiage the outliers in kPCA
reduced_kpca_2d = pd.DataFrame(data = reduced_kpca[:, 0:2], columns = ["PC1",
↪ "PC2"])
index = reduced_kpca_2d.index
outliers = reduced_kpca_2d[(reduced_kpca_2d["PC2"] > 0.2) &
↪ (reduced_kpca_2d["PC1"] < 0.5) ]



print("Num of outliers: ", len(outliers.index))
```

Num of outliers: 45

3.1 Visualizing the outliers

```
[48]: for outlier in outliers.index[1:4]:  
      visualize(G[outlier])
```





We can see that the outliers are associated with low/high number of nodes/edges.


3.2 Exploring Dataset properties

3.2.1 Node - Edge correlation

```
[9]: node_count = []
edge_count = []
for i in range(len(G)):
    node_count.append(len(G[i].nodes))
    edge_count.append(len(G[i].edges))
node_ser = pd.Series(node_count)
edge_ser = pd.Series(edge_count)
class_ser = pd.Series(classes)
df = pd.DataFrame({"class":class_ser,"nodes":node_ser,"edges":edge_ser})
```

```
[10]: plt.subplots(figsize=(16,8))
sns.scatterplot(data=df, x="nodes", y="edges", hue="class", palette="tab10")
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7fbf7f5c5b50>
```



As expected, we can see that a strong linear correlation between the nodes and edges. Also we can notice that in this vector representation the labels are more or less evenly distributed. Which makes sense because the enzymes size, connectiveness is not a major factor that determines the EC of an enzyme.

Let's plot visualize the node and edge distribution of our dataset.


3.2.2 Node/edge count distributions


```
[11]: # lets plot the node count distribution of the graphs
threshold_nodes= [10, 52]

node_lengths = [graph.number_of_nodes() for graph in G ]
plt.hist(node_lengths, bins = 30)
plt.axvline(x=threshold_nodes[0], color='r', linestyle='dashed', linewidth=2)
plt.axvline(x=threshold_nodes[1], color='r', linestyle='dashed', linewidth=2)
plt.xlabel("node count")
plt.ylabel("frequency")
plt.title("Node count distribution")
plt.show()

# lets plot the edge count distribution of the graphs
threshold_edges= [22, 100]

edge_lengths = [graph.number_of_edges() for graph in G ]
plt.hist(edge_lengths, bins = 30)
plt.axvline(x=threshold_edges[0], color='r', linestyle='dashed', linewidth=2)
plt.axvline(x=threshold_edges[1], color='r', linestyle='dashed', linewidth=2)
plt.xlabel("edge count")
plt.ylabel("frequency")
plt.title("Edge count distribution")
plt.show()
```





We can notice that there are quite some enzymes with extreme low/high node/edge counts.

For our preprocessing step we will apply edge and node count cutoffs, and remove the selected outliers because they are probably the source which make our dimensionality reduction less informative as they do not capture the essence of data i.e. characteristics of ground truth labels.

3.3 Preprocessing

```
[12]: # lets make some cutoffs to remove some extreme data points
G_filtered = []
indices_true = []
for graph, index in zip(G, range(0, len(G))):
    if threshold_nodes[0] < graph.number_of_nodes() < threshold_nodes[1] \
        and threshold_edges[0] < graph.number_of_edges() < threshold_edges[1] \
        and index not in outliers.index: #remove this line if you dont want
        ↪to filter out the outliers

        G_filtered.append(graph)
        indices_true.append(index)

indices_true = np.array(indices_true)
classes = dp.get_dataset(dataset)
```

```

classes = classes[indices_true] #subsetting the label list to remove OL
print("Size of filtered data: ",len(G_filtered))

```

Size of filtered data: 488

4 Dimensionality Reduction - WL

Here we want to try out clustering on different vector representations across different WL kernel iterations in order to see which WL iteration might capture the most amount of information. We also evaluate the clustering results with NMI metrics.

```

[ ]: # Here we use the WL kernel and for every WL iteration we apply DR and
      →Clustering and evaluate
      # the clustering methods
      # This is done on labeled and unlabeled dataset

for use_labels in [True, False]:
    if use_labels:
        base_path = os.path.join("kernels","node_labels")
        dataset = "ENZYMES"
    else:
        base_path = os.path.join("kernels","without_labels")
        dataset = "ENZYMES"

    print("#####")
    print("Load from ", base_path)
    print("#####")

    # dicts for storing nmi score
    nmis_kpca = {}
    nmis_tsvd = {}
    nmis_tsne = {}
    nmis_umap = {}
    nmis_spec = {}
    nmis_optics = {}

    nmis_kpca[dataset] = []
    nmis_tsvd[dataset] = []
    nmis_tsne[dataset] = []
    nmis_umap[dataset] = []
    nmis_spec[dataset] = []
    nmis_optics[dataset] = []

    # iterating over different WL iterations
    for iterations in range(1,6):
        print("#####")

```

```

print("Dataset: ", dataset)
print("Iteration: ", iterations)
print("#####")

#-----
# Load graph representation from the kernel
#-----

#Gram Matrix for the Weisfeiler-Lehman subtree kernel
gram = load_csv(os.path.
→join(base_path,f"{dataset}_gram_matrix_wl{iterations}.csv"))
gram = aux.normalize_gram_matrix(gram) # normalizing for dr
gram = gram[indices_true[:, None], indices_true] #removing outliers

#Sparse Vectors for the Weisfeiler-Lehmann subtree kernel
vec = load_sparse(os.path.
→join(base_path,f"{dataset}_vectors_wl{iterations}.npz"))
vec = vec[indices_true, :] #removing outliers
vec = csc_matrix(vec, dtype=np.int8).toarray()
#print(gram.shape, vec.shape)

#-----
# Dimensionality Reduction
#-----

# SVD
tsvd = TruncatedSVD(n_components=100)
reduced_tsvd = tsvd.fit_transform(vec)
#plot_dr(reduced_tsvd, "TSVD", classes)

# Kernel PCA
kpca = KernelPCA(n_components=100, kernel="precomputed")
reduced_kpca= kpca.fit_transform(gram)
#plot_dr(reduced_kpca, "Kernel PCA",classes)

# T-SNE
pca = TSNE(n_components=3)
reduced_tnse = pca.fit_transform(gram)
plot_dr(reduced_tnse, "T-SNE",classes)

# UMAP
reducer = umap.UMAP()
#scaled_gram = StandardScaler().fit_transform(gram)
#embedding = reducer.fit_transform(scaled_gram)
embedding = reducer.fit_transform(gram)

```

```

#plot_dr(embedding, "UMAP",classes)

k = len(set(classes.tolist()))
print("Number of labels: ", k)
d = {0:"TSVD",1:"KPCA", 2:"T-SNE", 3:"UMAP",}
n_d = {0:nmis_tsvd, 1:nmis_kpca, 2:nmis_tsne, 3:nmis_umap}

for i, rep_i in enumerate([reduced_tsvd, reduced_kpca, reduced_tnse,
↪embedding]):

    # KMeans
    kmeans = KMeans(n_clusters=10 ,random_state=0).fit(reduced_kpca)
    plot_dr(rep_i, "KMeans in: {0}".format(d[i]),kmeans.labels_)
    km_nmi = normalized_mutual_info_score(kmeans.labels_ , classes)
    n_d[i][dataset].append(km_nmi)
    print(f"KMeans NMI:{km_nmi:.4f}")

    #Spectral Clustering
    spec = SpectralClustering(n_clusters=k, affinity="precomputed")
    spec.fit(gram)
    spec_nmi = normalized_mutual_info_score(spec.labels_,classes)
    plot_dr(reduced_kpca, "Spectral Clustering",spec.labels_)
    nmis_spec[dataset].append(spec_nmi)
    print("#####")

    # OPTICS
    print("OPTICS")
    clust = OPTICS(min_samples=50, xi=.05, min_cluster_size=.01)
    clust.fit(gram)
    labels_050 = cluster_optics_dbscan(reachability=clust.reachability_,
                                       core_distances=clust.core_distances_,
                                       ordering=clust.ordering_, eps=0.5)

    space = np.arange(len(gram))
    reachability = clust.reachability_[clust.ordering_]
    labels = clust.labels_[clust.ordering_]
    optics_nmi = normalized_mutual_info_score(clust.labels_ , classes)
    nmis_optics[dataset].append(optics_nmi)

    plt.figure(figsize=(10, 7))
    G = gridspec.GridSpec(2, 3)
    ax1 = plt.subplot(G[0, :])
    ax2 = plt.subplot(G[1, 0])

    # Reachability plot

```

```

colors = ['g.', 'r.', 'b.', 'y.', 'c.']
for klass, color in zip(range(0, 5), colors):
    Xk = space[labels == klass]
    Rk = reachability[labels == klass]
    ax1.plot(Xk, Rk, color, alpha=0.3)
ax1.plot(space[labels == -1], reachability[labels == -1], 'k.', alpha=0.
↪3)

ax1.plot(space, np.full_like(space, 2., dtype=float), 'k-', alpha=0.5)
ax1.plot(space, np.full_like(space, 0.5, dtype=float), 'k-', alpha=0.5)
ax1.set_ylabel('Reachability (epsilon distance)')
ax1.set_title('Reachability Plot')

# OPTICS - clusters
colors = ['g.', 'r.', 'b.', 'y.', 'c.']
for klass, color in zip(range(0, 5), colors):
    Xk = gram[clust.labels_ == klass]
    ax2.plot(Xk[:, 0], Xk[:, 1], color, alpha=0.3)
ax2.plot(gram[clust.labels_ == -1, 0], gram[clust.labels_ == -1, 1],
↪'k+', alpha=0.1)
ax2.set_title('Automatic Clustering\nOPTICS')

res = {"KPCA":nmis_kpca,"TSVD": nmis_tsvd, "SPEC": nmis_spec,
       "OPTICS":nmis_optics, "TSNE":nmis_tsne, "UMAP": nmis_umap, }
for key, value in res.items():
    nmi_df = pd.DataFrame(res[key])
    nmi_df.to_csv(os.path.join(base_path, key + ".csv"))


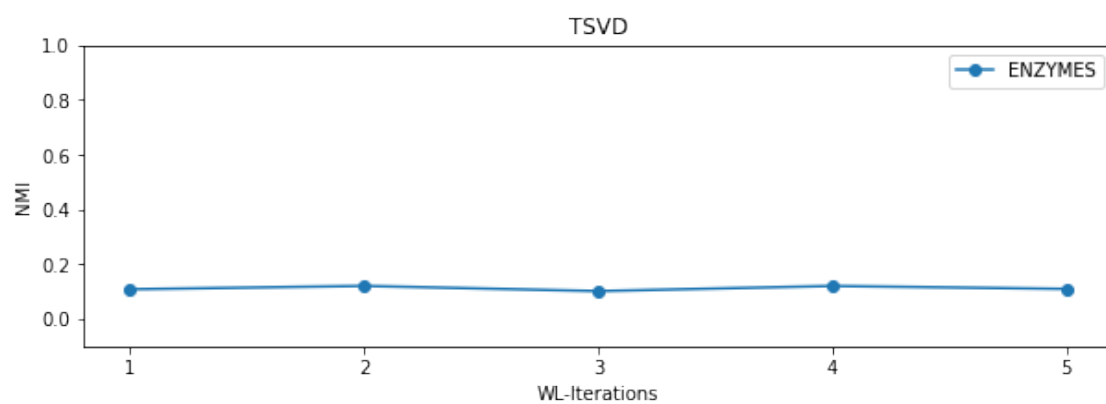
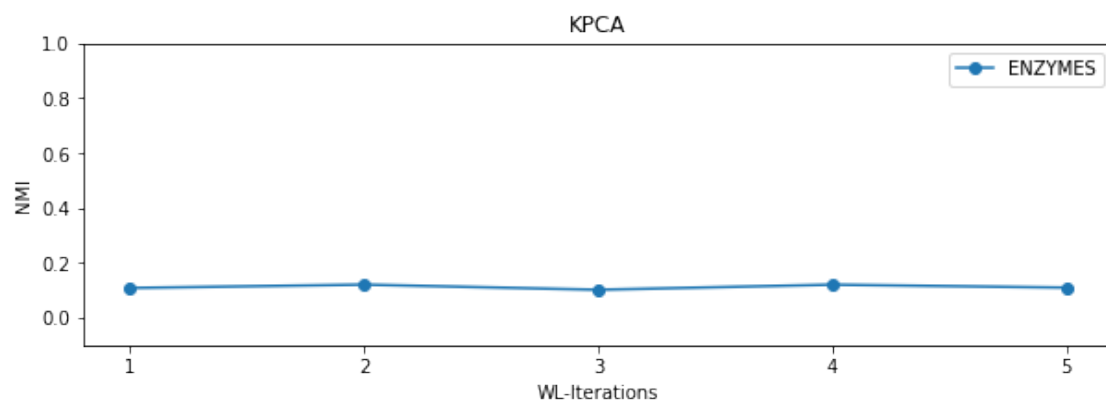
```

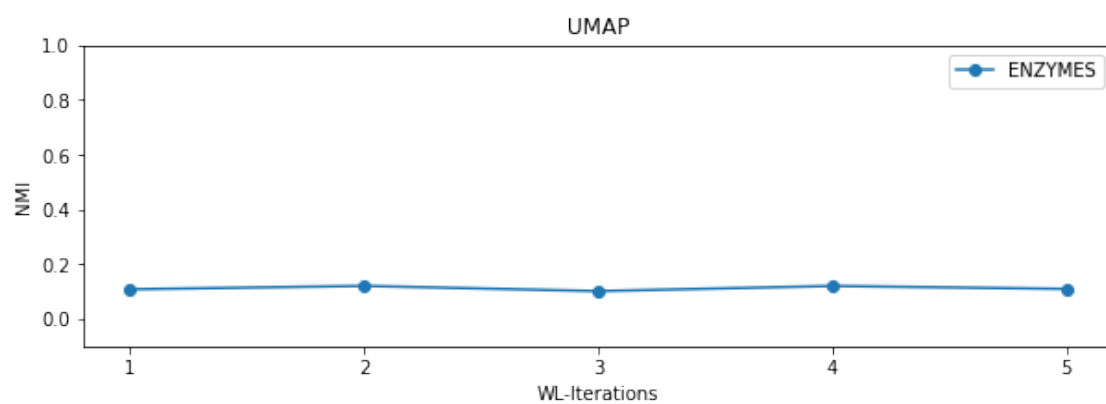
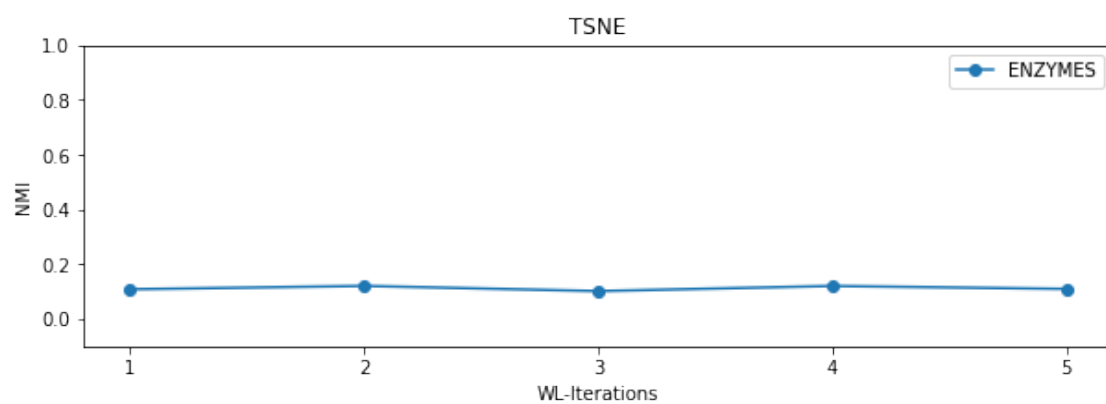
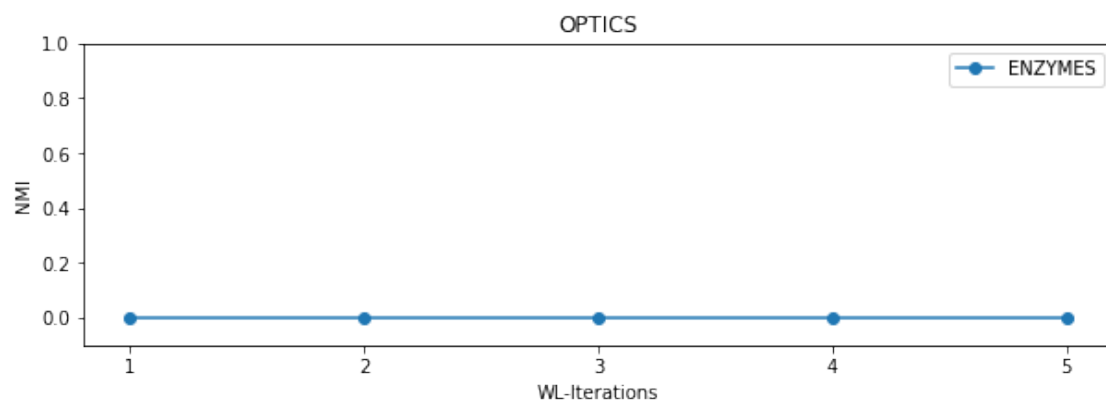
```

[16]: result_names = ["KPCA", "TSVD", "SPEC", "OPTICS", "TSNE", "UMAP"]
for version in ["node_labels", "without_labels"]:
    print(f"#####{version}#####")
    for name_i in result_names:
        path_i = os.path.join("kernels", version, name_i+".csv")
        nmi_df = pd.read_csv(path_i, index_col=0)
        fig, ax = plt.subplots(figsize=(10,3))
        ax.set_ylabel("NMI")
        ax.set_xlabel("WL-Iterations")
        ax.set_ylim([-0.1,1])
        ax.set_xticks([0,1,2,3,4])
        ax.set_xticklabels([1,2,3,4,5])
        ax.set_title(name_i)
        nmi_df.plot(marker="o", ax=ax)
        plt.show();

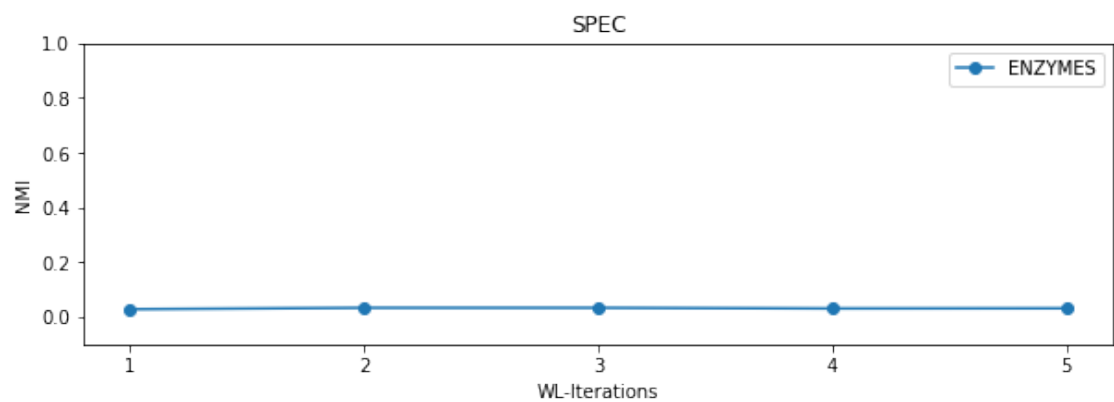
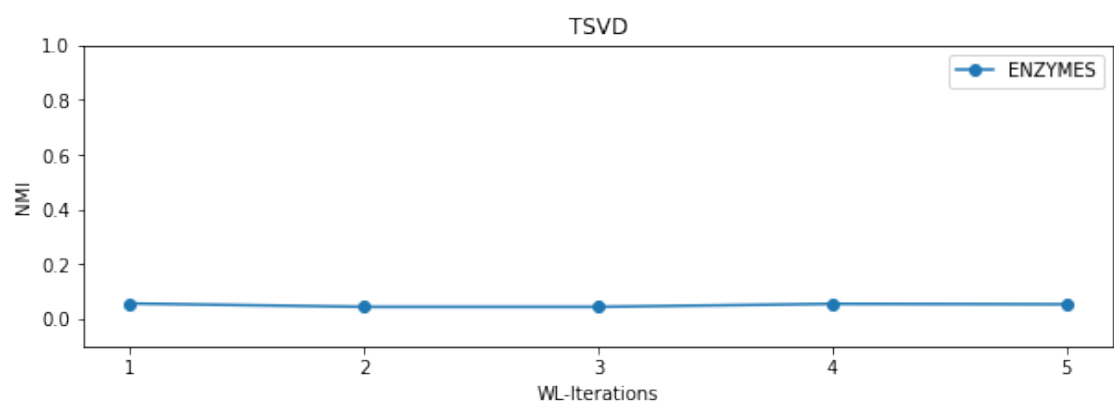
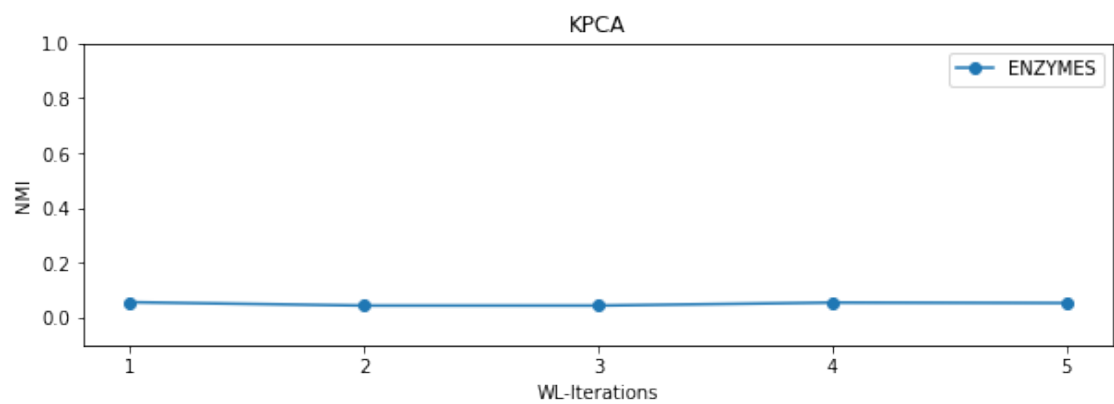
```

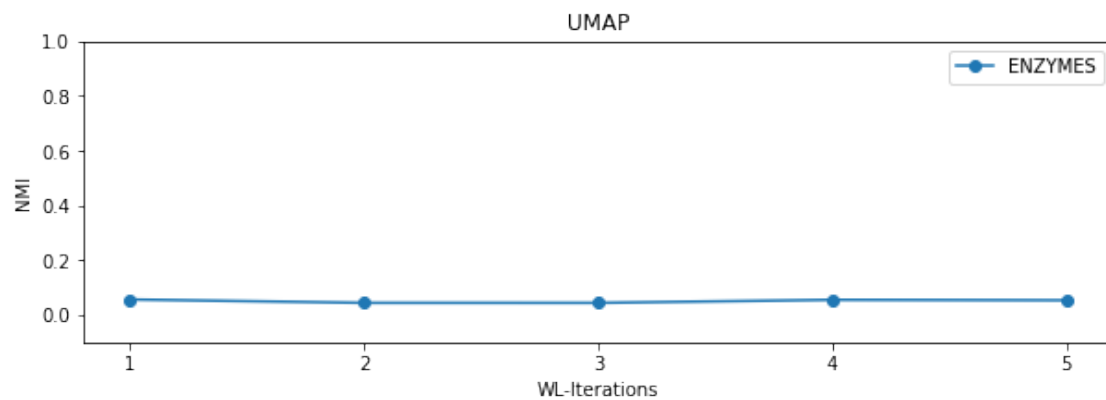
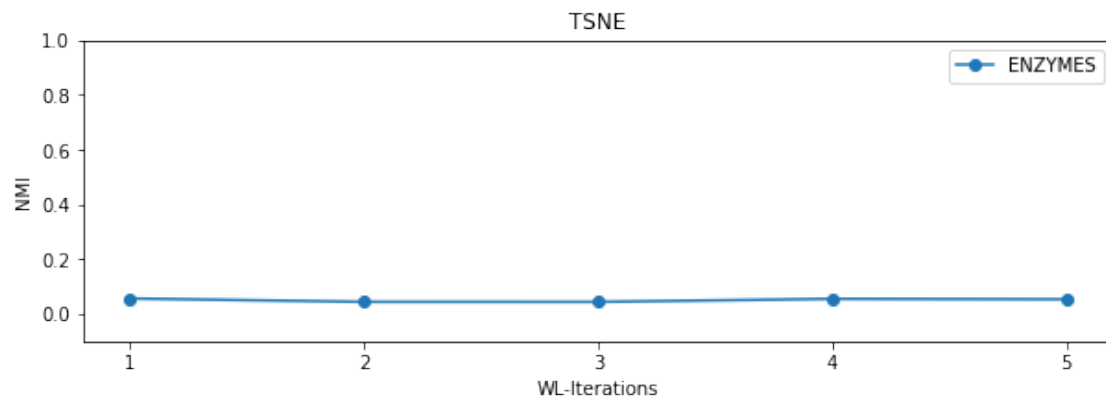
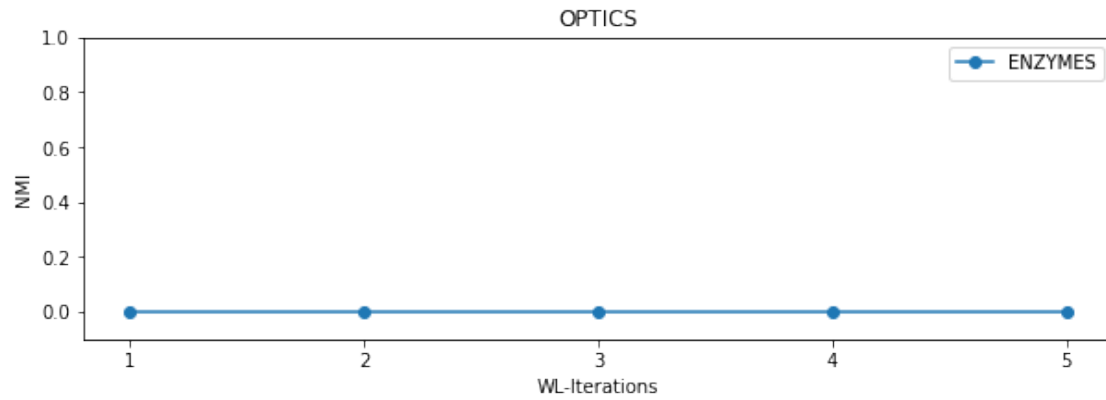
```
#####node_labels#####
```



#####without_labels#####





We can notice that the clustering results are a bit better in iteration 2 and 4 e.g. in kMeans clustering with UMAP, TSNE, kPCA. We will therefore use iteration 4 for our next part of the analysis.

Lets visualize the WL4 DR representations.

```
[17]: # trying different iterations
dataset = "ENZYMES"
iterations = 4
base_path = os.path.join("kernels", "node_labels")
#base_path = os.path.join("kernels", "without_labels")
classes = dp.get_dataset(dataset)
classes = classes[indices_true]
#-----
# Load graph representation from the kernel
#-----

#Gram Matrix for the Weisfeiler-Lehman subtree kernel
gram = load_csv(os.path.join(base_path, f"{dataset}_gram_matrix_wl{iterations}."
    ↪ "csv"))
gram = aux.normalize_gram_matrix(gram)
gram = gram[indices_true[:, None], indices_true] #removing outliers

#Sparse Vectors for the Weisfeiler-Lehman subtree kernel
vec = load_sparse(os.path.join(base_path, f"{dataset}_vectors_wl{iterations}."
    ↪ "npz"))
vec = vec[indices_true, :] #removing outliers
vec = csc_matrix(vec, dtype=np.int8).toarray()
print(gram.shape, vec.shape)

# SVD
tsvd = TruncatedSVD(n_components=50)
reduced_tsvd = tsvd.fit_transform(vec)
plot_dr(reduced_tsvd, "TSVD", classes)


# Kernel PCA
kpca = KernelPCA(n_components=100, kernel="precomputed")
reduced_kpca= kpca.fit_transform(gram)
plot_dr(reduced_kpca, "Kernel PCA", classes)


# T-SNE
pca = TSNE(n_components=3)
reduced_tsne = pca.fit_transform(gram)
plot_dr(reduced_tsne, "T-SNE", classes)


# UMAP
reducer = umap.UMAP()
scaled_gram = StandardScaler().fit_transform(gram)
embedding = reducer.fit_transform(scaled_gram)
```

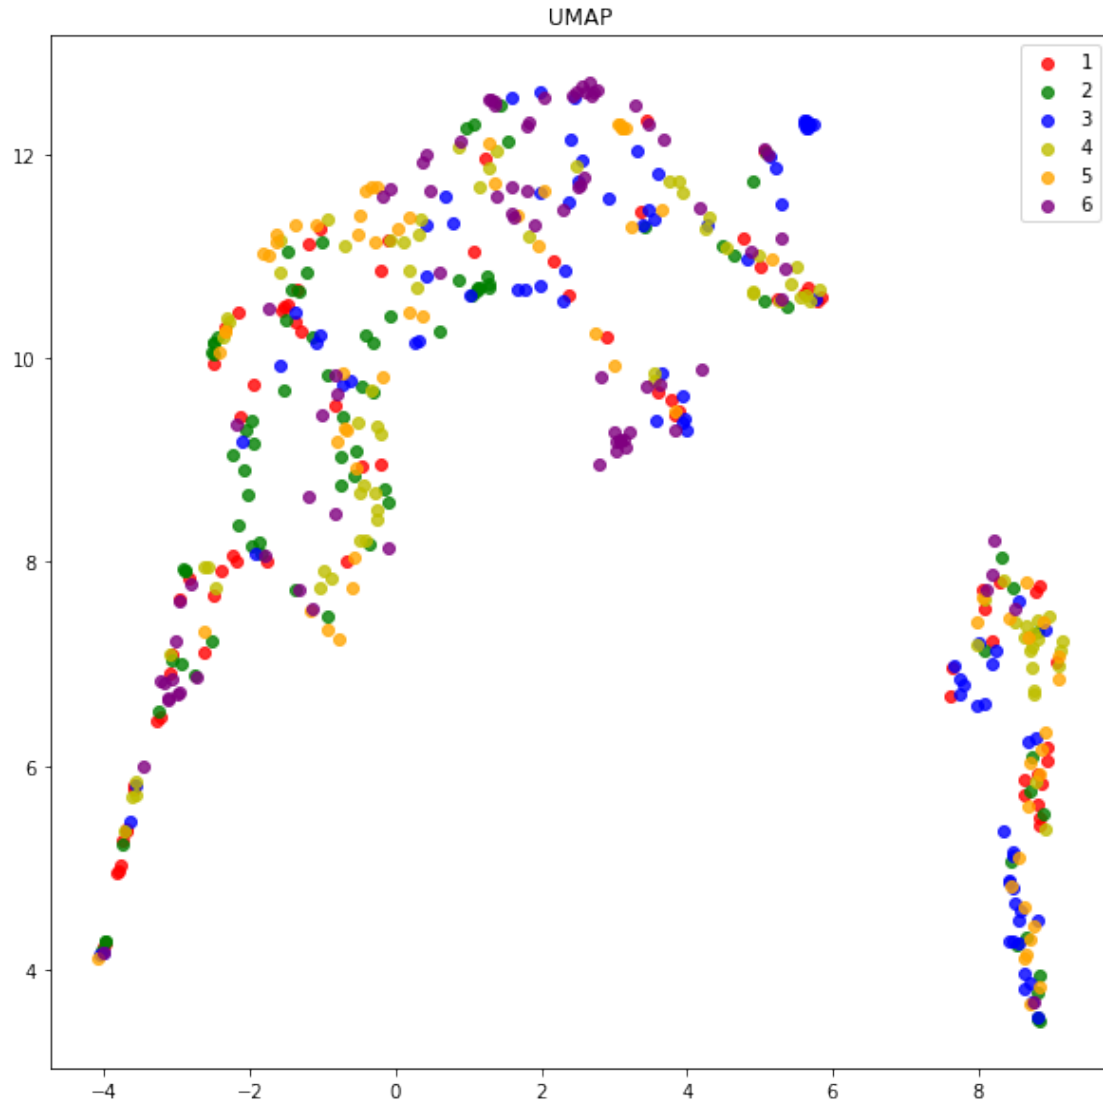
```
plot_dr(embedding, "UMAP", classes)
```

(488, 488) (488, 59599)









We can notice that there are some groupings of our truth labels in some places.

4.1 Clustering Evaluation

```
[18]: dataset = "ENZYMES"
iterations = 4
base_path = os.path.join("kernels", "node_labels")
#base_path = os.path.join("kernels", "without_labels")
classes = dp.get_dataset(dataset)
classes = classes[indices_true]

# dicts for storing nmi score
nmis_kpca = {}
```

```

nmis_tsvd = {}
nmis_tsne = {}
nmis_umap = {}
nmis_spec = {}
nmis_optics = {}
nmis_hier = {}
nmis_hisc = {}

nmis_kpca[dataset] = []
nmis_tsvd[dataset] = []
nmis_tsne[dataset] = []
nmis_umap[dataset] = []
nmis_spec[dataset] = []
nmis_optics[dataset] = []
nmis_hier[dataset] = []
nmis_hisc[dataset] = []

iterations = 4
print("#####")
print("Dataset: ", dataset)
print("Iteration: ", iterations)
print("#####")

#-----
# Load graph representation from the kernel
#-----

#Gram Matrix for the Weisfeiler-Lehman subtree kernel
gram = load_csv(os.path.join(base_path,f"{dataset}_gram_matrix_wl{iterations}."
    ↳csv"))
gram = aux.normalize_gram_matrix(gram) # normalizing for dr
gram = gram[indices_true[:, None], indices_true] #removing outliers

#Sparse Vectors for the Weisfeiler-Lehmann subtree kernel
vec = load_sparse(os.path.join(base_path,f"{dataset}_vectors_wl{iterations}."
    ↳npz"))
vec = vec[indices_true, :] #removing outliers
vec = csc_matrix(vec, dtype=np.int8).toarray()
print(gram.shape, vec.shape)

#-----
# Dimensionality Reduction
#-----

```

```

# SVD
tsvd = TruncatedSVD(n_components=100)
reduced_tsvd = tsvd.fit_transform(vec)
plot_dr(reduced_tsvd, "TSVD", classes)

# Kernel PCA
kpca = KernelPCA(n_components=100, kernel="precomputed")
reduced_kpca= kpca.fit_transform(gram)
plot_dr(reduced_kpca, "Kernel PCA", classes)

# T-SNE
pca = TSNE(n_components=3)
reduced_tnse = pca.fit_transform(gram)
plot_dr(reduced_tnse, "T-SNE", classes)

# UMAP
reducer = umap.UMAP()
#scaled_gram = StandardScaler().fit_transform(gram)
#embedding = reducer.fit_transform(scaled_gram)
embedding = reducer.fit_transform(gram)
plot_dr(embedding, "UMAP", classes)

k = len(set(classes.tolist()))
print("Number of labels: ", k)
d = {0:"TSVD", 1:"KPca", 2:"T-SNE", 3:"UMAP",}
n_d = {0:nmis_tsvd, 1:nmis_kpca, 2:nmis_tsne, 3:nmis_umap}

for i, rep_i in enumerate([reduced_tsvd, reduced_kpca, reduced_tnse,
    ↪embedding]):
    print(d[i])

    # KMeans
    print("KMeans")
    kmeans = KMeans(n_clusters=10 ,random_state=0).fit(rep_i)
    plot_dr(reduced_kpca, "KMeans in: {0}".format(d[i]), kmeans.labels_)
    km_nmi = normalized_mutual_info_score(kmeans.labels_ , classes)
    n_d[i][dataset].append(km_nmi)
    print(f"KMeans NMI:{km_nmi:.4f}")

# Hierarchical clustering
clustering = AgglomerativeClustering(n_clusters=10, linkage="ward").fit(gram)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)

```

```

nmis_hier[dataset].append(hc_nmi)
plot_dr(reduced_kpca, "hierarchical ward", clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

#Spectral Clustering
spec = SpectralClustering(n_clusters=k, affinity="precomputed")
spec.fit(gram)
spec_nmi = normalized_mutual_info_score(spec.labels_, classes)
plot_dr(reduced_kpca, "Spectral Clustering", clustering.labels_)
nmis_spec[dataset].append(spec_nmi)
print("#####")

# OPTICS
print("OPTICS")
clust = OPTICS(min_samples=10, xi=.001, min_cluster_size=.01)
clust.fit(gram)
labels_050 = cluster_optics_dbscan(reachability=clust.reachability_,
                                   core_distances=clust.core_distances_,
                                   ordering=clust.ordering_, eps=0.5)

space = np.arange(len(gram))
reachability = clust.reachability_[clust.ordering_]
labels = clust.labels_[clust.ordering_]
optics_nmi = normalized_mutual_info_score(clust.labels_, classes)
print("optics nmi: ", optics_nmi)
nmis_optics[dataset].append(optics_nmi)

plt.figure(figsize=(10, 7))
G = gridspec.GridSpec(2, 3)
ax1 = plt.subplot(G[0, :])
ax2 = plt.subplot(G[1, 0])

# Reachability plot
colors = ['g.', 'r.', 'b.', 'y.', 'c.']
for klass, color in zip(range(0, 5), colors):
    Xk = space[labels == klass]
    Rk = reachability[labels == klass]
    ax1.plot(Xk, Rk, color, alpha=0.3)
ax1.plot(space[labels == -1], reachability[labels == -1], 'k.', alpha=0.3)
ax1.plot(space, np.full_like(space, 2., dtype=float), 'k-', alpha=0.5)
ax1.plot(space, np.full_like(space, 0.5, dtype=float), 'k-', alpha=0.5)
ax1.set_ylabel('Reachability (epsilon distance)')
ax1.set_title('Reachability Plot')

# OPTICS - clusters

```

```

colors = ['g.', 'r.', 'b.', 'y.', 'c.']
for klass, color in zip(range(0, 5), colors):
    Xk = gram[clust.labels_ == klass]
    ax2.plot(Xk[:, 0], Xk[:, 1], color, alpha=0.3)
ax2.plot(gram[clust.labels_ == -1, 0], gram[clust.labels_ == -1, 1], 'k+', alpha=0.1)
ax2.set_title('Automatic Clustering\nOPTICS')
plt.show()

# self-implemented HiSC
print("HiSC")
gram = load_csv("./kernels/node_labels/ENZYMES_gram_matrix_wl4.csv")
gram = aux.normalize_gram_matrix(gram)
gram = gram[indices_true[:, None], indices_true]

label_file = load_csv("../tud_benchmark/datasets/ENZYMES/ENZYMES/raw/
    ENZYMES_graph_labels.txt")[indices_true]
label_file = label_file.astype(int)
np.savetxt("./kernels/node_labels/ENZYMES_gram_matrix_wl4_filtered.csv", gram,
    delimiter=";")
np.savetxt("../tud_benchmark/datasets/ENZYMES/ENZYMES/raw/
    ENZYMES_graph_labels_filtered.txt", label_file,
    fmt="%d", newline="\n")
dataset_file = "./kernels/node_labels/ENZYMES_gram_matrix_wl4_filtered.csv"
label_file = "../tud_benchmark/datasets/ENZYMES/ENZYMES/raw/
    ENZYMES_graph_labels_filtered.txt"

data, labels_true = process_csv(dataset_file, input_filename_labels=label_file,
    sep=";")

alpha = 0.007
k = 4
clus_ord = HiSC(data, alpha, k, verbose=False, elki=True)
threshold=0.8
dimensions=[(0,9),(1,11)]

labels = get_clusters(clus_ord, threshold=threshold)

hisc_nmi = normalized_mutual_info_score(labels, classes)
print(hisc_nmi)
nmis_hisc[dataset].append(hisc_nmi)

reachability_plot(data, clus_ord, labels, dimensions=dimensions)

```

```


res = {"KPCA":nmis_kpca,"TSVD": nmis_tsvd, "SPEC": nmis_spec,
      "OPTICS":nmis_optics, "TSNE":nmis_tsne, "UMAP": nmis_umap, "HIER":
      ↪nmis_hier,
      "HISC": nmis_hisc}
for key, value in res.items():
    nmi_df = pd.DataFrame(res[key])
    nmi_df.to_csv(os.path.join(base_path, key + ".csv"))


```


```


#####
Dataset:  ENZYMES
Iteration:  4
#####
(488, 488) (488, 59599)

```










Number of labels: 6

TSVD


KMeans



KMeans NMI:0.0923

KPCA


KMeans



KMeans NMI:0.1205

T-SNE


KMeans




KMeans NMI:0.0979

UMAP


KMeans




KMeans NMI:0.0968



H clustering NMI:0.0956





```
#####  
OPTICS  
optics nmi: 0.13358307709649445
```



HISC

Running HiSC, input dataset has 488 entries with 488 dimensions
 0.025576301418530626al structure | 100.0% complete






```
[19]: result_names = ["KPCA", "TSVD", "SPEC", "OPTICS", "TSNE", "UMAP", "HIER", "HISC"]
      ↪ "HISC"]
version = "node_labels"
#version = "without_labels"

nmis = []
print(f"#####{version}#####")
for name_i in result_names:
    path_i = os.path.join("kernels", version, name_i+".csv")

    nmi_df = pd.read_csv(path_i, index_col=0)
    nmis.append(float(nmi_df.values[0]))

fig = plt.figure()
ax = fig.add_axes([0,0,1,1])
ax.bar(result_names,nmis)
ax.set_ylabel('NMI')
plt.show()
```

```
#####node_labels#####
```




We notice that most of the clusters do not represent the ground truth labels with any clustering method. For the next part we will only use kmeans clustering as its NMI is descent. Also HiSC clustering shows that it is not suitable for this kind of data.

4.2 Interpreting DR and Clusters found in UMAP

```
[20]: costs = []
kmeans_runs = []
k_list = [1, 2, 5, 10, 15, 20, 25, 30, 35, 40]
for k in k_list:
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(embedding)
    costs.append(kmeans.inertia_)
    kmeans_runs.append(deepcopy(kmeans))
```

```
[21]: def scree_plot(k_list, costs):
    plt.figure(figsize=(10,5))
    plt.plot(k_list, costs, marker="o")
    plt.xlabel('Number of Clusters')
    # plt.ylabel('KMeans loss')
    plt.title('Search for the elbow to determine the optimal k')
    plt.show();
scree_plot(k_list, costs)
```




We use number of clusters between 5-10 as they are the most optimal


```
[22]: plot_dr(embedding, "UMAP", classes)
      plt.show();

      # Instantiate the clustering model
      kmeans = KMeans(n_clusters=8)
      kmeans.fit(reduced_kpca)
      print("KMeans")

      plot_dr(embedding, "kmeans", kmeans.labels_)
```



KMeans



4.2.1 Visualizing data points in found clusters

```
[23]: # select first 2 nearest neighbours in for each cluster
nr_nearest = 2
nearest_indices = []
for cluster_i in set(kmeans.labels_):
    mask = (kmeans.labels_ == cluster_i)
    selection = reduced_kpca[mask]
    print(f"number of data points in cluster {cluster_i}: {selection.shape[0]}")
    center_i = kmeans.cluster_centers_[cluster_i].reshape(1,-1)
    distances_i = pairwise_distances(center_i, reduced_kpca)
    nearest_indices.append(np.argsort(distances_i, ) [0] [0:nr_nearest])
    print(f"Nearest data points in cluster {cluster_i}: {nearest_indices[-1]}")
```

```

number of data points in cluster 0: 84
Nearest data points in cluster 0: [233 240]
number of data points in cluster 1: 32
Nearest data points in cluster 1: [455 319]
number of data points in cluster 2: 100
Nearest data points in cluster 2: [276 133]
number of data points in cluster 3: 78
Nearest data points in cluster 3: [ 66 333]
number of data points in cluster 4: 42
Nearest data points in cluster 4: [475 409]
number of data points in cluster 5: 58
Nearest data points in cluster 5: [ 72 192]
number of data points in cluster 6: 64
Nearest data points in cluster 6: [485 481]
number of data points in cluster 7: 30
Nearest data points in cluster 7: [463 462]


```


```

[24]: for i, indices in enumerate(nearest_indices):
      print(f"Cluster {i}, Indices {indices}")
      G_selected = select_from_list(G_filtered, indices)
      for g_i in G_selected:
          print(g_i)
          visualize(g_i)



```

Cluster 0, Indices [233 240]







Cluster 1, Indices [455 319]





Cluster 2, Indices [276 133]







Cluster 3, Indices [66 333]





Cluster 4, Indices [475 409]







Cluster 5, Indices [72 192]





Cluster 6, Indices [485 481]





Cluster 7, Indices [463 462]



We can see that the clusters are strongly driven by the graph structure e.g. degree, node/edge counts.

```
[26]: def plot_degree_dist(G, labels):
    degrees= []
    for graph, label in zip(G, labels):
        degrees = [graph.degree(n) for n in graph.nodes()]
        plt.hist(degrees, alpha=0.8, label=label, bins= 10)
    plt.legend()
    plt.show()

graph_centers = []
for i, indices in enumerate(nearest_indices):
    G_selected = select_from_list(G_filtered, indices)
    graph_centers.append(G_selected[0])



plot_degree_dist(graph_centers[0:3], labels = ["1", "2", "3"])
plot_degree_dist(graph_centers[3:6], labels = ["4", "5", "6"])



avg_degree= []
for graph in G_filtered:
    degrees = [graph.degree(n) for n in graph.nodes()]
    avg= sum(degrees)/len(degrees)
    avg_degree.append(avg)

avg_centrality= []
for graph in G_filtered:
    degrees = nx.degree_centrality(graph)
    avg= sum(degrees.values())/len(degrees)
    avg_centrality.append(avg)

ax = sns.violinplot(x=classes, y=avg_degree)
plt.title("Average Degree in EC classes")
plt.show()

ax = sns.violinplot(x=classes, y=avg_centrality)
plt.title("Average Degree centrality in EC classes")
plt.show()
```





We can see that the EC classes have different degree distributions

4.3 Analyzing the labels

```
[27]: G = tud_to_networkx(dataset)
class1_index = []
class2_index = []
class3_index = []
class4_index = []
class5_index = []
class6_index = []

for ind in classes:
    if ind == 1:
        class1_index.append(ind)

    if ind == 2:
        class2_index.append(ind)

    if ind == 3:
        class3_index.append(ind)

    if ind == 4:
        class4_index.append(ind)

    if ind == 5:
        class5_index.append(ind)

    if ind == 6:
        class6_index.append(ind)


ec_classes = [class1_index, class2_index, class3_index, class4_index,
→class5_index, class6_index]
```

```
[28]: labels = ["EC1", "EC2", "EC3", "EC4", "EC5", "EC6"]


for label, ec in zip(labels, ec_classes):
    print(label)

    visualize(G_filtered[ec[10]])
```


EC1




EC2




EC3




EC4



EC5



EC6



```
[29]: def get_label(G, attribute):
    data = []
    for graph in G:
        g_data = 0
        for node in range(0, len(graph)):
            g_data += (graph.nodes.data()[node][attribute][0])

        data.append(g_data/len(graph))
    return(data)

def get_property(G, attribute, index):
    data = []
    for gid, graph in enumerate(G):
        g_data = 0
        for node in range(0, len(graph)):
            g_data += graph.nodes.data()[node][attribute][index]
        data.append(g_data/len(graph))

    return(data)
```

```
def get_EC(G):
    data = []

    for i in range(0, len(classes)):
        data.append(classes[i])
    return data
```

```
[43]: data = {'Node_label': get_label(G_filtered, "labels"),
              'Node_count': [graph.number_of_nodes() for graph in G_filtered ],
              'Edge_count': [graph.number_of_edges() for graph in G_filtered ],
              'AA_length': get_property(G_filtered, "attributes", 0),
              'low Waals': get_property(G_filtered, "attributes", 1),
              'med Waals': get_property(G_filtered, "attributes", 2),
              'high Waals': get_property(G_filtered, "attributes", 3),
              'low Hydro': get_property(G_filtered, "attributes", 4),
              'med Hydro': get_property(G_filtered, "attributes", 5),
              'high Hydro': get_property(G_filtered, "attributes", 6),
              'low Polarity': get_property(G_filtered, "attributes", 7),
              'med Polarity': get_property(G_filtered, "attributes", 8),
              'high Polarity': get_property(G_filtered, "attributes", 9),
              'low Polariz': get_property(G_filtered, "attributes", 10),
              'med Polariz': get_property(G_filtered, "attributes", 11),
              'high Polariz': get_property(G_filtered, "attributes", 12),
              '3d length': get_property(G_filtered, "attributes", 13),
              'Total Waals': get_property(G_filtered, "attributes", 14),
              'Total Hydro': get_property(G_filtered, "attributes", 15),
              'Total Polarity': get_property(G_filtered, "attributes", 16),
              'Total Polariz': get_property(G_filtered, "attributes", 17),
              'Avg_degree': avg_degree, 'Avg_degree_centrality': avg_centrality,

              }

df = pd.DataFrame(data)
df.head()
```

```
[43]:
```

	Node_label	Node_count	Edge_count	AA_length	low Waals	med Waals	\
0	1.351351	37	84	7.756757	12.753363	28.328919	
1	1.347826	23	51	9.434783	15.357012	35.580870	
2	1.240000	25	46	9.240000	14.165496	33.696000	
3	1.250000	24	45	8.875000	13.915350	32.902083	
4	1.217391	23	45	9.913043	15.458343	36.126957	

	high Waals	low Hydro	med Hydro	high Hydro	...	low Polariz	\
0	0.921081	1.299595	64.213514	2.189189	...	2.405405	

1	1.813913	1.626304	79.643478	2.217391	...	2.565217
2	0.684400	1.541600	78.580000	2.600000	...	2.800000
3	0.568750	1.501792	75.883333	2.416667	...	2.583333
4	0.715217	1.651348	84.521739	2.826087	...	3.000000



	med Polariz	high Polariz	3d length	Total Waals	Total Hydro	\
0	2.756757	3.027027	2.027027	2.702703	1.864865	
1	3.391304	3.913043	1.739130	3.782609	1.913043	
2	3.160000	3.560000	2.080000	3.600000	2.400000	
3	3.000000	3.416667	1.875000	3.583333	2.250000	
4	3.391304	3.826087	2.217391	3.869565	2.608696	



	Total Polarity	Total Polariz	Avg_degree	Avg_degree_centrality
0	3.729730	2.162162	4.540541	0.126126
1	4.521739	3.000000	4.434783	0.201581
2	4.360000	2.480000	3.680000	0.153333
3	4.208333	2.416667	3.750000	0.163043
4	4.608696	2.695652	3.913043	0.177866

[5 rows x 23 columns]


4.4 Interpreting UMAP with Featureplots of node attributes

```
[44]: for feature in df.columns:
        sns.scatterplot(data=df, x=np.array(embedding[:,0]), y=np.array(embedding[:,1]), hue=feature, )
        plt.title("Feature plot: {}".format(feature))
        plt.show()
```







Feature plot: low Waals




Feature plot: med Waals




Feature plot: high Waals




Feature plot: low Hydro





Feature plot: med Hydro




Feature plot: high Hydro










Feature plot: high Polarity







Feature plot: low Polariz














We can notice that the Node_label has a strong effect on our UMAP1 vector. Also Node count and edge count have a strong effect on our UMAP2 vector.

```
[32]: x = df.values #returns a numpy array
      #df = preprocessing.normalize(x)
      min_max_scaler = preprocessing.MinMaxScaler()
      x_scaled = min_max_scaler.fit_transform(x)
      df = pd.DataFrame(x_scaled, columns= df.columns)

      data_norm= StandardScaler().fit_transform(x)

      df.head()
```

```
[32]:
```

	Node_label	AA_length	low Waals	med Waals	high Waals	low Hydro	\
0	0.392687	0.178778	0.373598	0.476537	0.246369	0.465153	
1	0.388747	0.270306	0.449869	0.598526	0.338841	0.582090	
2	0.268235	0.259682	0.414965	0.566819	0.221855	0.551772	
3	0.279412	0.239773	0.407637	0.553465	0.209877	0.537524	
4	0.242967	0.296393	0.452838	0.607712	0.225047	0.591053	

	med Hydro	high Hydro	low Polarity	med Polarity	...	low Polariz	\
0	0.535866	0.382577	0.493090	0.483840	...	0.373655	
1	0.664630	0.387505	0.610663	0.671329	...	0.398480	
2	0.655755	0.454369	0.602353	0.554965	...	0.434951	
3	0.633252	0.422330	0.585219	0.540793	...	0.401294	

4	0.705340	0.493879	0.635845	0.603223	...	0.466019
---	----------	----------	----------	----------	-----	----------

	med Polariz	high Polariz	3d length	Total Waals	Total Hydro	\
0	0.325521	0.328355	0.386100	0.470035	0.349662	
1	0.400449	0.424466	0.331263	0.657845	0.358696	
2	0.373137	0.386169	0.396190	0.626087	0.450000	
3	0.354244	0.370621	0.357143	0.623188	0.421875	
4	0.400449	0.415033	0.422360	0.672968	0.489130	

	Total Polarity	Total Polariz	Avg_degree	Avg_degree_centrality
0	0.491158	0.483840	0.742962	0.156686
1	0.595455	0.671329	0.703578	0.362102
2	0.574156	0.554965	0.422500	0.230754
3	0.554184	0.540793	0.448568	0.257189
4	0.606906	0.603223	0.509284	0.297540

[5 rows x 21 columns]

4.5 Investigating EC classes based on node features


```
[45]: # taking average values for each EC class
df.insert(1, "EC", get_EC(classes), True)
df = df.drop(columns=["Avg_degree", "Avg_degree_centrality", "Node_count",
    ↪ "Edge_count"])
df_avg = df.groupby("EC").mean()
df_avg = df_avg.T
df_avg.head()
```

```
[45]: EC          1          2          3          4          5          6
Node_label  1.445366  1.446252  1.549962  1.510326  1.518212  1.457125
AA_length   7.720490  8.492352  7.485750  8.073502  7.801927  8.345070
low Waals   13.128959 14.967227 13.691902 14.707855 14.281851 11.869869
med Waals   25.073924 20.569776 25.489137 26.839942 25.592492 28.956227
high Waals   0.885641  0.662938  0.880228  0.961177  0.971233  0.936135
```

```
[34]: df_diff = df_avg.sub(df_avg[1], axis= "rows")
```

```
[35]: sns.heatmap(df_diff, cmap="coolwarm" )
plt.title("Difference of average EC node properties compared to EC1")
```


```
[35]: Text(0.5, 1.0, 'Difference of average EC node properties compared to EC1')
```

We can see that node attributes play also an important role in EC classification.

```
[36]: sns.heatmap(df_avg.corr(method='spearman'), cmap="Blues")
plt.title("Correlation heatmap of EC classes")
```

```
[36]: Text(0.5, 1.0, 'Correlation heatmap of EC classes')
```



We can see that EC 5,6 are highly correlated thus they probably share some similarities Also EC 1,3 are highly correlated

```
[37]: df = df.drop(columns=["EC"])
      df.head()
```

```
[37]:
```

	Node_label	AA_length	low Waals	med Waals	high Waals	low Hydro \
0	0.392687	0.178778	0.373598	0.476537	0.246369	0.465153
1	0.388747	0.270306	0.449869	0.598526	0.338841	0.582090
2	0.268235	0.259682	0.414965	0.566819	0.221855	0.551772
3	0.279412	0.239773	0.407637	0.553465	0.209877	0.537524
4	0.242967	0.296393	0.452838	0.607712	0.225047	0.591053

	med Hydro	high Hydro	low Polarity	med Polarity	high Polarity \
0	0.535866	0.382577	0.493090	0.483840	0.465696
1	0.664630	0.387505	0.610663	0.671329	0.624303
2	0.655755	0.454369	0.602353	0.554965	0.588718
3	0.633252	0.422330	0.585219	0.540793	0.590812
4	0.705340	0.493879	0.635845	0.603223	0.632107

	low Polariz	med Polariz	high Polariz	3d length	Total Waals \
0	0.373655	0.325521	0.328355	0.386100	0.470035
1	0.398480	0.400449	0.424466	0.331263	0.657845
2	0.434951	0.373137	0.386169	0.396190	0.626087

3	0.401294	0.354244	0.370621	0.357143	0.623188
4	0.466019	0.400449	0.415033	0.422360	0.672968


	Total Hydro	Total Polarity	Total Polariz
0	0.349662	0.491158	0.483840
1	0.358696	0.595455	0.671329
2	0.450000	0.574156	0.554965
3	0.421875	0.554184	0.540793
4	0.489130	0.606906	0.603223

```
[38]: labels = ["EC1", "EC2", "EC3", "EC4", "EC5", "EC6"]


for label, ec in zip(labels, ec_classes):
    print(label)

    visualize(G_filtered[ec[10]])
```


EC1




EC2



EC3



EC4




EC5


There is no clear structural difference between the EC classes

4.6 Additional: Quick analysis of Node attributes


```
[81]: sns.set_theme(style="whitegrid")
      ax = sns.boxplot(x=classes, y="Total Hydro", data=df)
```




```
[82]: sns.set_theme(style="whitegrid")
      ax = sns.boxplot(x=classes, y="Total Polariz", data=df)
```




```
[83]: sns.set_theme(style="whitegrid")
ax = sns.boxplot(x=classes, y="Total Waals", data=df)
```




```
[84]: sns.set_theme(style="whitegrid")
ax = sns.boxplot(x=classes, y="3d length", data=df)
```



```
[85]: sns.set_theme(style="whitegrid")
ax = sns.boxplot(x=classes, y="high Waals", data=df)
```





4.7 Dimensionality reduction based on Node attributes


```
[87]: import umap
      # Kernel PCA
      kpca = PCA(n_components=18)
      reduced_kpca= kpca.fit_transform(data_norm)
      plot_dr(reduced_kpca, "kPCA",classes)

      # T-SNE
      pca = TSNE(n_components=3)
      reduced_tsne = pca.fit_transform(df.values)
      plot_dr(reduced_tsne, "T-SNE",classes)

      # UMAP
      reducer = umap.UMAP()
      embedding = reducer.fit_transform(df.values)
      plot_dr(embedding, "UMAP",classes)
```







Here we can see that node attributes also have big importance in the EC classification. E.g. EC2 and EC6 are nicely separated.

5 Support Vector machine prediction

```
[88]: #sum on all iterations
for iterations in range(1,6,1):
    gram = load_csv(os.path.
    ↪join(base_path,f"{dataset}_gram_matrix_wl{iterations}.csv"))
    gram = aux.normalize_gram_matrix(gram)
    gram = gram[indices_true[:, None], indices_true]
```

```

# Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(gram, classes,
→test_size=0.2,random_state=109) # 70% training and 30% test
clf = svm.SVC(kernel='poly') #Create a svm Classifier, 'poly' kernel is the
→best by far!
clf.fit(X_train, y_train) #Train the model using the training sets
y_pred = clf.predict(X_test) #Predict the response for test dataset
print(f"Iteration{iterations} accuracy:",metrics.accuracy_score(y_test,
→y_pred)) # Model Accuracy: how often is the classifier correct?

```

```

Iteration1 accuracy: 0.4489795918367347
Iteration2 accuracy: 0.5612244897959183
Iteration3 accuracy: 0.5714285714285714
Iteration4 accuracy: 0.6224489795918368
Iteration5 accuracy: 0.6122448979591837

```

Best accuracy is reached on iterations 3-4, and significantly higher than what could result from randomness (~0.16).

```

[89]: #sum on all iterations without "preprocessing"
classes_real = dp.get_dataset(dataset)
for iterations in range(1,6,1):
    gram = load_csv(os.path.
→join(base_path,f"{dataset}_gram_matrix_wl{iterations}.csv"))
    gram = aux.normalize_gram_matrix(gram)

    # Split dataset into training set and test set
    X_train, X_test, y_train, y_test = train_test_split(gram, classes_real,
→test_size=0.2,random_state=109) # 70% training and 30% test
    clf = svm.SVC(kernel='poly') #Create a svm Classifier, 'poly' kernel is the
→best by far!
    clf.fit(X_train, y_train) #Train the model using the training sets
    y_pred = clf.predict(X_test) #Predict the response for test dataset
    print(f"Iteration{iterations} accuracy:",metrics.accuracy_score(y_test,
→y_pred)) # Model Accuracy: how often is the classifier correct?

```

```

Iteration1 accuracy: 0.4583333333333333
Iteration2 accuracy: 0.575
Iteration3 accuracy: 0.625
Iteration4 accuracy: 0.625
Iteration5 accuracy: 0.6083333333333333

```

The svm works a bit better on the unprocessed kernel for most iterations, except for iteration5, but the difference is small.

6 Trying different kernels out

6.1 Graphlet kernel

6.1.1 Dim red

```
[37]: use_edge_labels = False
for use_labels in [True]:
    if use_labels:
        base_path = os.path.join("kernels", "node_labels")
        datasets = ["ENZYMES"]
    else:
        base_path = os.path.join("kernels", "without_labels")
        datasets = ["ENZYMES"]
    print("Load from ", base_path)
    nmis_kpca = {}
    nmis_tsvd = {}
    nmis_tsne = {}
    nmis_spec = {}
    dataset = datasets[0]

    nmis_kpca[dataset] = []
    nmis_tsvd[dataset] = []
    nmis_tsne[dataset] = []
    nmis_spec[dataset] = []

    print("#####")
    print("Dataset ", dataset)
    #print("Iteration ", iterations)
    print("#####")

    #-----
    # Load graph representation from the kernel
    #-----
    #Gram Matrix for the Weisfeiler-Lehman subtree kernel
    gram = load_csv(os.path.join(base_path, f"{dataset}_gram_matrix_graphlet.
↪csv"))
    #gram = load_csv(os.path.
↪join(base_path, f"{dataset}_gram_matrix_shortestpath.csv"))
    gram = aux.normalize_gram_matrix(gram)
    gram = gram[indices_true[:, None], indices_true]

    #Sparse Vectors for the Weisfeiler-Lehman subtree kernel
    vec = load_sparse(os.path.join(base_path, f"{dataset}_vectors_graphlet.npz"))
    vec = vec[indices_true, :]
    #vec = load_sparse(os.path.join(base_path, f"{dataset}_vectors_shortestpath.
↪npz"))
    print(gram.shape, vec.shape)
```

```

#-----
# Dimensionality Reduction
#-----

# SVD
tsvd = TruncatedSVD(n_components=50)
reduced_tsvd = tsvd.fit_transform(vec)
plot_dr(reduced_tsvd, "TSVD", classes)

# Kernel PCA
kpca = KernelPCA(n_components=50, kernel="precomputed")
reduced_kpca= kpca.fit_transform(gram)
plot_dr(reduced_kpca, "Kernel PCA", classes)

# T-SNE
pca = TSNE(n_components=3)
reduced_tnse = pca.fit_transform(gram)
plot_dr(reduced_tnse, "T-SNE", classes)


k = len(set(classes.tolist()))
print("n_clusters: ", k)
d = {0:"TSVD", 1:"KPCA", 2:"T-SNE"}
n_d = {0:nmis_tsvd, 1:nmis_kpca, 2:nmis_tsne}


```


```

Load from kernels\node_labels
#####
Dataset ENZYMES
#####
(488, 488) (488, 3455)

```





n_clusters: 6

6.1.2 Clustering

```
[38]: print("Dimensionality reduction: kPCA")
      plot_dr(reduced_kpca, "kPCA", classes)
      plt.show();

      # KMeans
      print("KMeans")
      kmeans = KMeans(n_clusters=10, random_state=0).fit(reduced_kpca)
      km_nmi = normalized_mutual_info_score(kmeans.labels_, classes)
      print(len(kmeans.labels_))
      print(len(classes))
```

```

plot_dr(reduced_kpca, "kmeans", kmeans.labels_)
print(f"KMeans NMI:{km_nmi:.4f}")

# Hierarchical clustering
clustering = AgglomerativeClustering(n_clusters=30, linkage="ward").
    ↳fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical ward", clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

clustering = AgglomerativeClustering(n_clusters=30, linkage="complete").
    ↳fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical complete", clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")


clustering = AgglomerativeClustering(n_clusters=10, linkage="average").
    ↳fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical average", clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

clustering = AgglomerativeClustering(n_clusters=10, linkage="single").
    ↳fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical single", clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

# Apply Subkmeans
k = len(set(classes.tolist()))
print("SubKMeans")
nrkm = NrKmeans(n_clusters=[k,1])#, allow_larger_noise_space=False)
nrkm.fit(reduced_kpca, best_of_n_rounds=10, verbose=False)
print("Found Cluster Subspaces: ", nrkm.m)
subkm_nmi = normalized_mutual_info_score(nrkm.labels[0], classes)
print(f"Subkmeans NMI:{subkm_nmi:.4f}")
#n_d[i][dataset].append(subkm_nmi)
# Plot rotated space
V = nrkm.V
rotated = np.dot(reduced_kpca,V)
reduced_df = pd.DataFrame(rotated[:,0:2])
reduced_df["labels"] = nrkm.labels[0] #classes[true_nodes]
sns.pairplot(reduced_df, hue="labels", diag_kind="hist", palette=("tab10"))
plt.show();

```


Dimensionality reduction: kPCA




KMeans

488


488




KMeans NMI:0.0828




H clustering NMI:0.1710



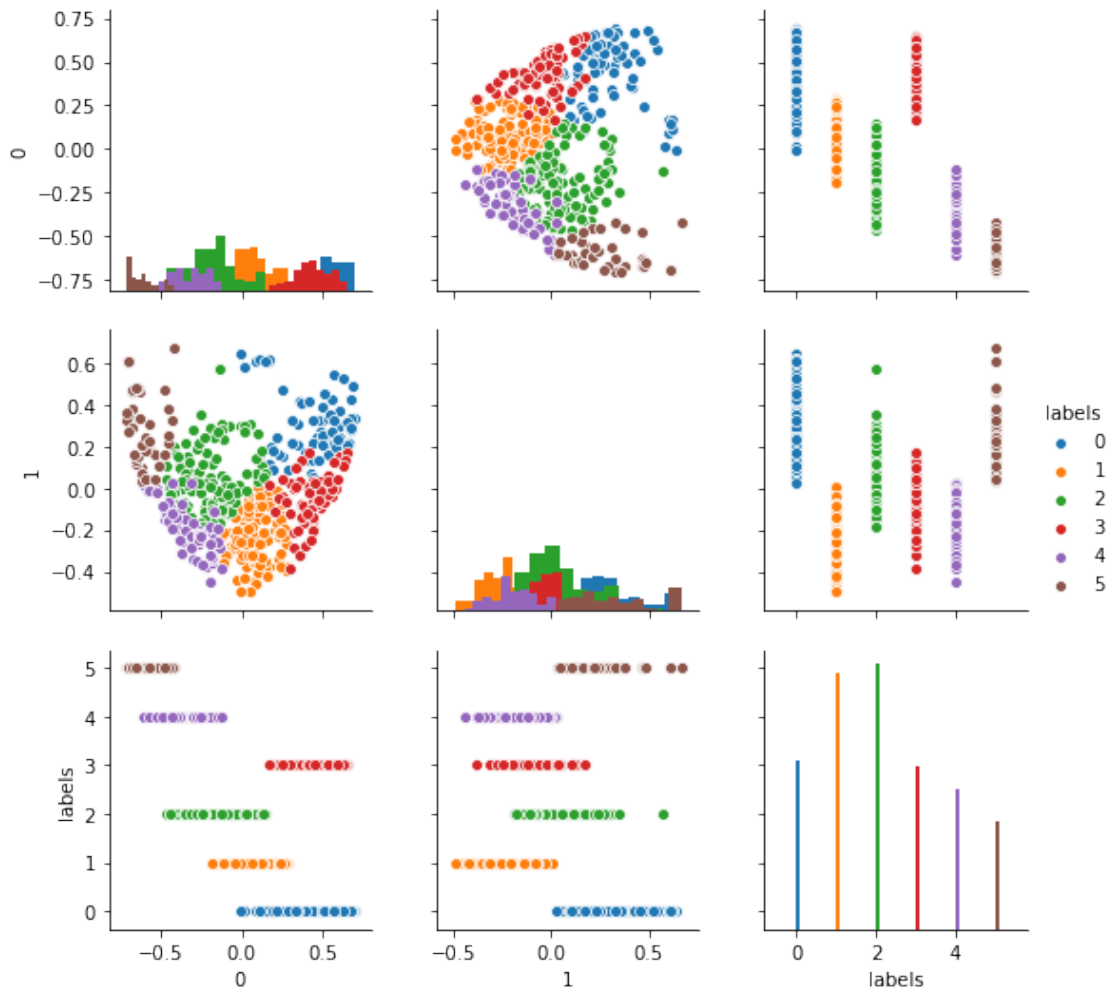
H clustering NMI:0.1576



H clustering NMI:0.0887



H clustering NMI:0.0551
SubKMeans
Found Cluster Subspaces: [5, 45]
Subkmeans NMI:0.0717



6.1.3 SVM

```
[39]: # Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(gram, classes, test_size=0.
    ↪ 2, random_state=109) # 70% training and 30% test
clf = svm.SVC(kernel='poly') # Create a svm Classifier, 'poly' kernel is the
    ↪ best by far!
clf.fit(X_train, y_train) # Train the model using the training sets
y_pred = clf.predict(X_test) # Predict the response for test dataset
print("Graphlet accuracy:", metrics.accuracy_score(y_test, y_pred)) # Model
    ↪ Accuracy: how often is the classifier correct?
```

Graphlet accuracy: 0.32653061224489793

6.2 Shortest path kernel

6.2.1 Dim red

```
[40]: use_edge_labels = False
for use_labels in [True]:
    if use_labels:
        base_path = os.path.join("kernels", "node_labels")
        datasets = ["ENZYMES"]
    else:
        base_path = os.path.join("kernels", "without_labels")
        datasets = ["ENZYMES"]
    print("Load from ", base_path)
    nmis_kpca = {}
    nmis_tsvd = {}
    nmis_tsne = {}
    nmis_spec = {}
    dataset = datasets[0]

    nmis_kpca[dataset] = []
    nmis_tsvd[dataset] = []
    nmis_tsne[dataset] = []
    nmis_spec[dataset] = []

    print("#####")
    print("Dataset ", dataset)
    #print("Iteration ", iterations)
    print("#####")

    #-----
    # Load graph representation from the kernel
    #-----
    #Gram Matrix for the Weisfeiler-Lehman subtree kernel
    gram = load_csv(os.path.join(base_path, f"{dataset}_gram_matrix_shortestpath.
→csv"))
    #gram = load_csv(os.path.
→join(base_path, f"{dataset}_gram_matrix_shortestpath.csv"))
    gram = aux.normalize_gram_matrix(gram)
    gram = gram[indices_true[:, None], indices_true]

    #Sparse Vectors for the Weisfeiler-Lehmann subtree kernel
    vec = load_sparse(os.path.join(base_path, f"{dataset}_vectors_shortestpath.
→npz"))
    #vec = load_sparse(os.path.join(base_path, f"{dataset}_vectors_shortestpath.
→npz"))
    vec = vec[indices_true, :]
    print(gram.shape, vec.shape)
```

```

#-----
# Dimensionality Reduction
#-----

# SVD
tsvd = TruncatedSVD(n_components=50)
reduced_tsvd = tsvd.fit_transform(vec)
plot_dr(reduced_tsvd, "TSVD", classes)

# Kernel PCA
kpca = KernelPCA(n_components=50, kernel="precomputed")
reduced_kpca= kpca.fit_transform(gram)
plot_dr(reduced_kpca, "Kernel PCA",classes)

# T-SNE
pca = TSNE(n_components=3)
reduced_tnse = pca.fit_transform(gram)
plot_dr(reduced_tnse, "T-SNE",classes)


k = len(set(classes.tolist()))
print("n_clusters: ", k)
d = {0:"TSVD",1:"KPCA", 2:"T-SNE"}
n_d = {0:nmis_tsvd, 1:nmis_kpca, 2:nmis_tsne}


```


```

Load from kernels\node_labels
#####
Dataset ENZYMES
#####
(488, 488) (488, 25460)

```







n_clusters: 6

6.2.2 Clustering

```
[41]: print("Dimensionality reduction: kPCA")
      plot_dr(reduced_kpca, "kPCA", classes)
      plt.show();

      # KMeans
      print("KMeans")
      kmeans = KMeans(n_clusters=10, random_state=0).fit(reduced_kpca)
      km_nmi = normalized_mutual_info_score(kmeans.labels_, classes)
      print(len(kmeans.labels_))
      print(len(classes))
```



```

plot_dr(reduced_kpca, "kmeans",kmeans.labels_)
print(f"KMeans NMI:{km_nmi:.4f}")

# Hierarchical clustering
clustering = AgglomerativeClustering(n_clusters=30, linkage="ward").
    ↪fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical ward",clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

clustering = AgglomerativeClustering(n_clusters=30, linkage="complete").
    ↪fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical complete",clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")


clustering = AgglomerativeClustering(n_clusters=10, linkage="average").
    ↪fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical average",clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

clustering = AgglomerativeClustering(n_clusters=10, linkage="single").
    ↪fit(reduced_kpca)
hc_nmi = normalized_mutual_info_score(clustering.labels_ , classes)
plot_dr(reduced_kpca, "hierarchical single",clustering.labels_)
print(f"H clustering NMI:{hc_nmi:.4f}")

# Apply Subkmeans
k = len(set(classes.tolist()))
print("SubKMeans")
nrkm = NrKmeans(n_clusters=[k,1])#, allow_larger_noise_space=False)
nrkm.fit(reduced_kpca, best_of_n_rounds=10, verbose=False)
print("Found Cluster Subspaces: ", nrkm.m)
subkm_nmi = normalized_mutual_info_score(nrkm.labels[0], classes)
print(f"Subkmeans NMI:{subkm_nmi:.4f}")
#n_d[i][dataset].append(subkm_nmi)
# Plot rotated space
V = nrkm.V
rotated = np.dot(reduced_kpca,V)
reduced_df = pd.DataFrame(rotated[:,0:2])
reduced_df["labels"] = nrkm.labels[0] #classes[true_nodes]
sns.pairplot(reduced_df, hue="labels", diag_kind="hist", palette=("tab10"))
plt.show();

```


Dimensionality reduction: kPCA




KMeans

488


488




KMeans NMI:0.1091




H clustering NMI:0.1866



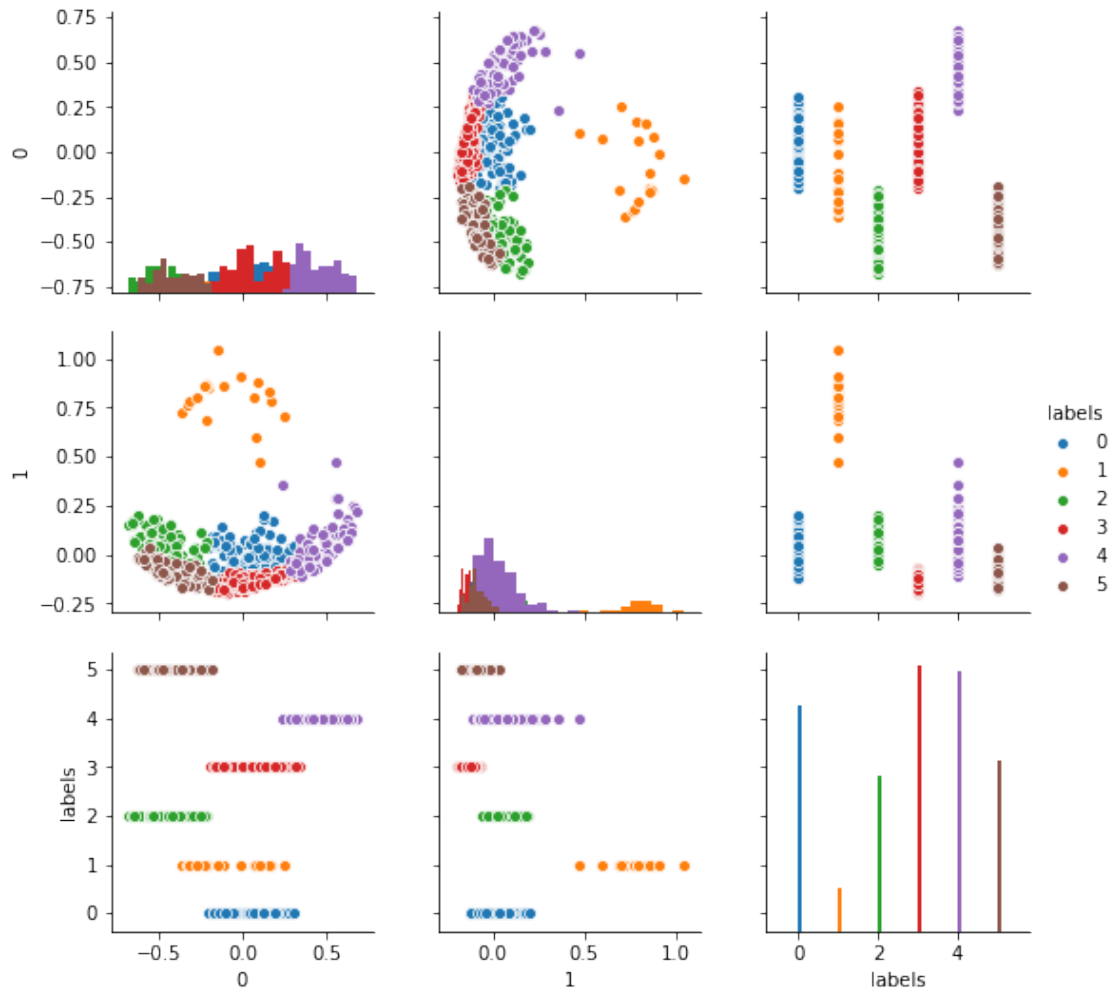
H clustering NMI:0.1601



H clustering NMI:0.0602



H clustering NMI:0.0508
SubKMeans
Found Cluster Subspaces: [5, 45]
Subkmeans NMI:0.0754



6.2.3 SVM

```
[42]: # Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(gram, classes, test_size=0.
    ↪ 2, random_state=109) # 70% training and 30% test
clf = svm.SVC(kernel='poly') # Create a svm Classifier, 'poly' kernel is the
    ↪ best by far!
clf.fit(X_train, y_train) # Train the model using the training sets
y_pred = clf.predict(X_test) # Predict the response for test dataset
print("Shortest path accuracy:", metrics.accuracy_score(y_test, y_pred)) # Model
    ↪ Accuracy: how often is the classifier correct?
```

Shortest path accuracy: 0.4897959183673469

7 Conclusion

- UMAP vector representation is strongly driven by node labels (SSE elements), edge/node count
- Clusters do not represent the ground truth labels
- Clusters are formed based on different SSE composition, and structural differences e.g. degree distribution
- Node attributes play also an important role in EC classification which could be used for building better kernels